

CSCI 315: Data Structures

Performance Analysis

Paul E. West, PhD

Department of Computer Science
Charleston Southern University

August 29, 2019

Analysis of Algorithms

- Dilemma: you have two (or more) methods to solve problem, how to choose the BEST?
- One approach: implement each algorithm in C, test how long each takes to run.
- Problems:
 - Different implementations may cause an algorithm to run faster/slower
 - Some algorithms run faster on some computers
 - Algorithms may perform differently depending on data (e.g., sorting often depends on what is being sorted)

Better Approach Step 1

- characterize performance in terms of key operation(s)
- Sorting:
 - count number of times two values compared
 - count number of times two values swapped
- Search:
 - count number of times value being searched for is compared to values in array
- Recursive function:
 - count number of recursive calls

Better Approach Step 2

- Want to comment on the “general” performance of the algorithm
- Empirical: Measure for several examples, but what does this tell us in general?
- Analytical:
 - Instead, assess performance in an abstract manner
 - Idea: analyze performance as size of problem grows
 - Examples:
 - Sorting: how many comparisons for array of size N ?
 - Searching: #comparisons for array of size N
 - May be difficult to discover a reasonable formula

Analysis With Varying Results

- Example: for some sorting algorithms, a sorting routine may require as few as $N-1$ comparisons and as many as $\frac{N^2}{2}$
- Types of analyses:
 - Best-case: what is the fastest an algorithm can run for a problem of size N ?
 - Average-case: on average how fast does an algorithm run for a problem of size N ?
 - Worst-case: what is the longest an algorithm can run for a problem of size N ?
- Computer scientists *usually* use worst-case analysis

Notice: We Are **Estimating**

- What is often done is to approximate or estimate the performance of an algorithm
- Estimation is an important skill to learn and to use
- Example Question: How many hotdogs tall is the Empire State Building?

Simple Example

- Simpler Question: How tall is the Empire State Building?
- Answer: The ESB is 1250 feet tall.
- Assuming that a hotdog is 6 inches from end to end, you would need, $1250 * 2 = 2500$ hotdogs.



Analysis

- An objective way to evaluate the cost of an algorithm or code section.
- The cost is computed in terms of space or time, usually
- The goal is to have a meaningful way to compare algorithms based on a common measure.
- Complexity analysis has two phases,
 - Algorithm analysis
 - Complexity analysis

Algorithm Analysis

- Algorithm analysis requires a set of rules to determine how operations are to be counted.
- There is no generally accepted set of rules for algorithm analysis.
- In some cases, an exact count of operations is desired; in other cases, a general approximation is sufficient.
- The rules presented that follow are typical of those intended to produce an exact count of operations.

Rules

- 1 We assume an arbitrary time unit.
- 2 Execution of one of the following operations takes time 1:
 - 1 assignment operation
 - 2 single I/O operations
 - 3 single Boolean operations, numeric comparisons
 - 4 single arithmetic operations
 - 5 function return
 - 6 array index operations, pointer dereferences

More Rules

- 3 Running time of a selection statement (if, switch) is the time for the condition evaluation + the maximum of the running times for the individual clauses in the selection.
- 4 Loop execution time is the sum, over the number of times the loop is executed, of the body time + time for the loop check and update operations, + time for the loop setup.
- 5 Always assume that the loop executes the maximum number of iterations possible Running time of a function call is 1 for setup + the time for any parameter calculations + the time required for the execution of the function body.

Example 1

```
count = count + 1; // Cost: c1  
sum = sum + count; // Cost: c2
```

Total Cost = $c1 + c2$.

Since we assume '+' cost 1 and assignment cost 1, the total cost is 4.

Example 2

```
if (n < 0) { // Cost: c1
    absval = -n; // Cost: c2
} else {
    absval = n; // Cost: c3
}
```

- Total Cost $\leq c1 + \max(c2, c3)$
- $c1$ is the cost of boolean evaluation. Since there is 1 evaluation ($<$), $\text{Cost}(c1) = 1$.
- $c2$ is the cost of negating a number (1) + the cost of assignment (1). $\text{Cost}(c2) = 2$.
- $c3$ is the cost of assignment(1). $\text{Cost}(c3) = 1$
- Cost of the worse-case is 3.
- Cost of the best-case is 2.
- Average case is 2.5.

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1, \text{Cost}(c2) = 1, \text{Cost}(c3) = 1.$

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$.
- $\text{Cost}(c4) = 1 + 1 = 2$ (remember assignment and + both cost 1!).

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$.
- $\text{Cost}(c4) = 1 + 1 = 2$ (remember assignment and + both cost 1!).
- $\text{Cost}(c5) = 2$.

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$.
- $\text{Cost}(c4) = 1 + 1 = 2$ (remember assignment and + both cost 1!).
- $\text{Cost}(c5) = 2$.
- How many time does the loop execute?

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$.
- $\text{Cost}(c4) = 1 + 1 = 2$ (remember assignment and + both cost 1!).
- $\text{Cost}(c5) = 2$.
- How many time does the loop execute?
- Loop: n times, so total cost is:

Example 3

```
i = 1; // Cost: c1
sum = 0; // Cost: c2
while (i <= n) { // Cost: c3
    i = i + 1; // Cost: c4
    sum = sum + i; // Cost: c5
}
```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$.
- $\text{Cost}(c4) = 1 + 1 = 2$ (remember assignment and + both cost 1!).
- $\text{Cost}(c5) = 2$.
- How many time does the loop execute?
- Loop: n times, so total cost is:
- $\text{Total Cost} = c1 + c2 + (n+1)*c3 + n*c4 + n*c5 = c1 + c2 + c3 + n(c3 + c4 + c5)$

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

- Cost(c1) = 1, Cost(c2) = 1, Cost(c3) = 1, Cost(c4) = 1,
Cost(c5) = 1, Cost(c6) = 2, Cost(c7) = 2, Cost(c8) = 2

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

- Cost(c1) = 1, Cost(c2) = 1, Cost(c3) = 1, Cost(c4) = 1, Cost(c5) = 1, Cost(c6) = 2, Cost(c7) = 2, Cost(c8) = 2
- First (outer) while loop execution: n

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

- $\text{Cost}(c1) = 1$, $\text{Cost}(c2) = 1$, $\text{Cost}(c3) = 1$, $\text{Cost}(c4) = 1$,
 $\text{Cost}(c5) = 1$, $\text{Cost}(c6) = 2$, $\text{Cost}(c7) = 2$, $\text{Cost}(c8) = 2$
- First (outer) while loop execution: n
- Second (inner) while loop execution: n , total cost is:

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

- $\text{Cost}(c1) = 1, \text{Cost}(c2) = 1, \text{Cost}(c3) = 1, \text{Cost}(c4) = 1,$
 $\text{Cost}(c5) = 1, \text{Cost}(c6) = 2, \text{Cost}(c7) = 2, \text{Cost}(c8) = 2$
- First (outer) while loop execution: n
- Second (inner) while loop execution: n , total cost is:
- $$c1 + c2 + (n+1) * c3 + n * c4 + n * (n+1) * c5 + n * n * c6 + n * n * c7 + n * c8 =$$

$$c1 + c2 + c3 + n * (c3 + c4 + c8) + n * n * c5 + n * c5 + n * n * c6 + n * n * c7 =$$

$$c1 + c2 + c3 + n * (c3 + c4 + c5 + c8) + n * n(c5 + c6 + c7)$$

Nested Example

```

i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j = 1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1; // Cost c8
}

```

- $\text{Cost}(c1) = 1, \text{Cost}(c2) = 1, \text{Cost}(c3) = 1, \text{Cost}(c4) = 1,$
 $\text{Cost}(c5) = 1, \text{Cost}(c6) = 2, \text{Cost}(c7) = 2, \text{Cost}(c8) = 2$
- First (outer) while loop execution: n
- Second (inner) while loop execution: n , total cost is:
- $$c1 + c2 + (n+1) * c3 + n * c4 + n * (n+1) * c5 + n * n * c6 + n * n * c7 + n * c8 =$$

$$c1 + c2 + c3 + n * (c3 + c4 + c8) + n * n * c5 + n * c5 + n * n * c6 + n * n * c7 =$$

$$c1 + c2 + c3 + n * (c3 + c4 + c5 + c8) + n * n(c5 + c6 + c7)$$
- **Important Note:** $n*n$ (n^2) is the highest (largest) term!

Comparing Algorithms

- We measure an algorithm's time requirement as a function of the problem size.
- Problem size depends on the application: e.g. number of elements in a list for a sorting algorithm, the number disks for towers of hanoi.
- So, for instance, we say that (if the problem size is n)
 - Algorithm A requires $5 * n^2$ time units to solve a problem of size n .
 - Algorithm B requires $7 * n$ time units to solve a problem of size n .
- An algorithm's proportional time requirement is known as growth rate.
- We can compare the efficiency of two algorithms by comparing their growth rates.

Example

- Which is better?
 - $50N^2 + 31N^3 + 24N + 15$
 - $3N^2 + N + 21 + 4 * 3^N$

Example

- Which is better?

- $50N^2 + 31N^3 + 24N + 15$
- $3N^2 + N + 21 + 4 * 3^N$

- Well, it depends on N:

| N | $50N^2 + 31N^3 + 24N + 15$ | $3N^2 + N + 21 + 4 * 3^N$ |
|----|----------------------------|---------------------------|
| 1 | 120 | 37 |
| 2 | 511 | 71 |
| 3 | 1374 | 159 |
| 4 | 2895 | 397 |
| 5 | 5260 | 1073 |
| 6 | 8655 | 3051 |
| 7 | 13266 | 8923 |
| 8 | 19279 | 26465 |
| 9 | 26880 | 79005 |
| 10 | 36255 | 236527 |

What happened?

| N | $3N^2 + N + 21 + 4 * 3^N$ | $4 * 3^N$ | % of Total |
|----|---------------------------|-----------|------------|
| 1 | 37 | 12 | 32.4 |
| 2 | 71 | 36 | 50.7 |
| 3 | 159 | 108 | 67.9 |
| 4 | 397 | 324 | 81.6 |
| 5 | 1073 | 972 | 90.6 |
| 6 | 3051 | 2916 | 95.6 |
| 7 | 8923 | 8748 | 98.0 |
| 8 | 26465 | 26244 | 99.2 |
| 9 | 79005 | 78732 | 99.7 |
| 10 | 236527 | 236196 | 99.9 |

- One term dominated the others.
- This implies we *really* only care about the dominating (highest/largest) term.

As N Grows, Some Terms Dominate

| Function | N=10 | N=100 | N=1000 | N=10000 | N = 100000 |
|----------------|--------|-----------|------------|-------------|--------------|
| $\log_2 N$ | 3 | 6 | 9 | 13 | 16 |
| N | 10 | 100 | 1000 | 10000 | 100000 |
| $N * \log_2 N$ | 30 | 664 | 9965 | 10^5 | 10^6 |
| N^2 | 10^2 | 10^4 | 10^6 | 10^8 | 10^{10} |
| N^3 | 10^3 | 10^6 | 10^9 | 10^{12} | 10^{15} |
| 2^N | 10^3 | 10^{30} | 10^{301} | 10^{3010} | 10^{30103} |

- Ordering:

$$1 < \log_2 N < N < N * \log_2 N < N^2 < N^3 < 2^N < 3^N$$

Big O

- If Algorithm A requires time proportional to $f(n)$, Algorithm A is said to be order $f(n)$, and it is denoted as $O(f(n))$.
- The function $f(n)$ is called the algorithm's growth-rate function.
- Since the capital O is used in the notation, this notation is called the Big O notation.
- If Algorithm A requires time proportional to n^2 , it is $O(n^2)$.
- If Algorithm A requires time proportional to n , it is $O(n)$.

Example 1

- If an algorithm requires $n^2 - 3 * n + 10$ seconds to solve a problem size n . If constants k and n_0 exist such that $k * n^2 + n_0 > n^2 - 3 * n + 10$ for all n and n_0 .
- Then the algorithm is order n^2 (In fact, k is 3 and n_0 is 2)
- Thus, the algorithm requires no more than $k * n^2$ time units.
- So it is $O(n^2)$

Examples

- This is actually not that difficult. It is a game of “spot the highest term!”
- $50N^2 + 31N^3 + 24N + 15 = O(N^3)$
- $3N^2 + N + 21 + 4 * 3^N = O(3^N)$
- It can get somewhat tricky:
- $N(3 + N(9 + N)) + N^2 = O(N^3)$
- $N(10 + \log_2 N) + N = O(N * \log_2 N)$

Growth Rate Function Explained

- $O(1)$ Time requirement is constant, and it is independent of the problem's size.
- $O(\log_2 n)$ Time requirement for a logarithmic algorithm increases slowly as the problem size increases.
- $O(n)$ Time requirement for a linear algorithm increases directly with the size of the problem.
- $O(n * \log_2 n)$ Time requirement for a $n * \log_2 n$ algorithm increases more rapidly than a linear algorithm.
- $O(n^2)$ Time requirement for a quadratic algorithm increases rapidly with the size of the problem.
- $O(n^3)$ Time requirement for a cubic algorithm increases more rapidly with the size of the problem than the time requirement for a quadratic algorithm.
- $O(2^n)$ As the size of the problem increases, the time requirement is too rapid to be practical.

Example 1

```
i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    i = i + 1; // Cost c4
    sum = sum + i; // Cost c5
}
```

Example 1

```
i = 1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    i = i + 1; // Cost c4
    sum = sum + i; // Cost c5
}
```

- $$\begin{aligned} T(n) &= c1 + c2 + (n+1)*c3 + n*c4 + n*c5 \\ &= (c3+c4+c5)*n + (c1+c2+c3) \\ &= a*n + b \end{aligned}$$

- So, the growth-rate function for this algorithm is $O(n)$

Example 2

```
i=1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j=1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1 // Cost c8
}
```

Example 2

```

i=1; // Cost c1
sum = 0; // Cost c2
while (i <= n) { // Cost c3
    j=1; // Cost c4
    while (j <= n) { // Cost c5
        sum = sum + i; // Cost c6
        j = j + 1; // Cost c7
    }
    i = i + 1 // Cost c8
}

```

- $$\begin{aligned}
 T(n) &= c1 + c2 + (n+1)*c3 + n*c4 + \\
 &\quad n*(n+1)*c5 + n*n*c6 + n*n*c7 + n*c8 \\
 &= (c5+c6+c7)*n*n + (c3+c4+c5+c8)*n + (c1+c2+c3) \\
 &= a*n*n + b*n + c
 \end{aligned}$$
- So, the growth-rate function for this algorithm is $O(n^2)$

Example 3

```
for (i=1; i<=n; i++) { // Cost(c1)
  for (j=1; j<=i; j++) { // Cost(c2)
    for (k=1; k<=j; k++) { // Cost(c3)
      x=x+1; // Cost(c4)
    }
  }
}
```


Example 3

```

for (i=1; i<=n; i++) { // Cost(c1)
  for (j=1; j<=i; j++) { // Cost(c2)
    for (k=1; k<=j; k++) { // Cost(c3)
      x=x+1; // Cost(c4)
    }
  }
}

```

- $T(n) = c1*(n+1) + c2*((n+1)*(n+2)) / 2 + c3* (\text{ estimated: } (n * (n + 1) * (2n + 1)) / 6) + c4*(\text{ estimated: } (n * (n + 1) * (2n + 1)) / 6)$
 $= a*n^3 + b*n^2 + c*n + d$
- So, the growth-rate function for this algorithm is $O(n^3)$
- Notice:** You do NOT need to know the exact number of iterations to find Big-O.

Example 4

Unfortunately, recursive can be hard...

```
void hanoi(int n, char source, char dest, char spare) { // Cost of function call
    if (n > 0) { // Cost(c1)
        hanoi(n-1, source, spare, dest); // Cost(c2)
        cout << "Move top disk from pole " << source
            << " to pole " << dest << endl; // Cost(c3)
        hanoi(n-1, spare, dest, source); // Cost(c4)
    } }
```

- By now, I hope you see that constance costs are virtually surpfulous when working with Big O.
- To find the growth-rate function for a recursive algorithm, we have to solve its recurrence relation.
- You will learn how to do this in Discrete Structures.

Example 4 continued

- What is the cost of $\text{hanoi}(n, 'A', 'B', 'C')$?
- when $n=0$ $T(0) = c_1$
- when $n>0$ $T(n) = c_1 + c_2 + T(n-1) + c_3 + c_4 + T(n-1)$
 $= 2 * T(n-1) + (c_1 + c_2 + c_3 + c_4)$
 $= 2 * T(n-1) + c \rightarrow$ recurrence equation for the growth-rate function of hanoi-towers algorithm
- Now, we have to solve this recurrence equation to find the growth-rate function of hanoi-towers algorithm
- This turns out to be $O(2^n)$ because for every N we make $2(n-1)$ calls.

Crash Course

- gnuplot makes graphs
- type “gnuplot” at your terminal
- type “plot sin(x) with line”
- type “plot sin(x) with point”
- Type “set terminal postscript color”
- Type “set output “nameofplot.ps” ”
- Type “replot” or “rep”

Continued

- Type “set title “plotname” ”
- Type “set ylabel “ylabel” ”
- Type “set xlabel “xlabel” ”
- to covert to something readable (like pdf). On the command line do:
\$ ps2pdf nameofplot.ps
- Then you should have nameofplot.pdf
\$ evince nameofplot.pdf

Huge Time Saver!

- The commands to gnuplot can be saved to a file and then automatically used:
- simple.plot:

```
set terminal postscript color
set output "simple.ps"
set ylabel "time (seconds)"
set xlabel "size"
```

...

- \$ gnuplot simple.plot
- \$ ps2pdf simple.ps
- \$ evince simple.pdf
- Will display the graph.
- This is especially useful if you put it in a makefile!

```
plot:
    gnuplot simple.plot
    ps2pdf simple.ps
    evince simple.pdf
```

Perf

- Tool to find what part of the code is running slow.
- Steps:
 - 1 compile with debug flags! (-g)
 - 2 run:
\$ sudo perf record <your-programs-name>
 - 3 view the report:
\$ sudo perf report
- Now you should see what functions are taking up all the time!

```
Samples: 991 of event 'cycles:ppp', Event count (approx.): 2061166086
Overhead Command Shared Object Symbol
95.29% a.out a.out [.] sumOfOneTo
4.39% a.out ld-2.28.so [.] 0x000000000000a25a
0.20% a.out a.out [.] sumOfOneToSquared
0.08% a.out [kernel.kallsyms] [k] __do_page_fault
0.03% a.out [kernel.kallsyms] [k] kmem_cache_alloc
0.00% perf_4. [kernel.kallsyms] [k] native_sched_clock
0.00% perf_4. [kernel.kallsyms] [k] native_apic_mem_write
0.00% perf_4. [kernel.kallsyms] [k] native_write_msr
```

Emperical Measurement

- While Analytical measurement of performance is important sometimes an emperical approach is most useful.
- The naive way to take timings:
 - 1 Record time as start
 - 2 Run section of code you wish to time
 - 3 Record time as end
 - 4 You answer is (end - start).
- While there are numerous issues with this approach, it will give sufficient approximate timings for this class.


```

#include <iostream> // To print
#include <time.h> // Required for taking timings

int main(int argc, char *argv[]) { // Standard main heading.
    /* clock_t is the data type for storing timing information.
     * We must make two variables, one for the start and the other to capture
     * the difference.
     */
    clock_t start, diff;

    // timeAmount is used to print out the time in seconds.
    double timeAmount;

    // We want to run our algorithm over varying sizes.
    for (int i = 1000; i < 1000000; i += 1000) {
        // Capture the start clock
        start = clock();

        // This is where your algorithm should be called.
        functionCallToYourAlgorithm(i);

        // Capture the clock and subtract the start to get the total time elapsed.
        diff = clock() - start;

        // Convert clock_t into seconds as a floating point number.
        timeAmount = diff * 1.0 / CLOCKS_PER_SEC;

        // Print out first the size (i) and then the elapsed time.
        std::cout << i << " " << timeAmount << "\n";

        // We flush to ensure the timings is printed out.
        std::cout << std::flush;
    }
}

```

Example Output

```

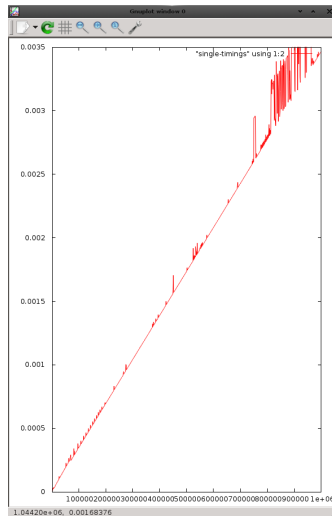
1000 4e-06
2000 8e-06
3000 1.2e-05
4000 2.5e-05
5000 2.9e-05
6000 2.4e-05
7000 3.5e-05
8000 2.9e-05
9000 3.2e-05
10000 3.5e-05
11000 3.9e-05
12000 4.2e-05
13000 4.5e-05
14000 5.2e-05
15000 5.6e-05
16000 6e-05
17000 6.5e-05
18000 6.8e-05
19000 7e-05
20000 7.6e-05

```

- First is the size (1000) and second is the number of seconds (pretty small.)

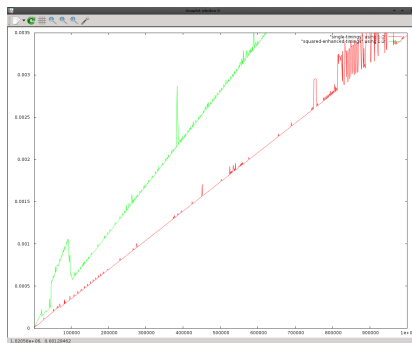
Plotting

- First lets us assume we have the full listing in a file named 'single-timings'
- With gnuplot we can simply graph the timings:
plot [:][:] "single-timings" using 1:2 with line



Multiple Data Collections

- In addition to 'single-timings' let us assume we have another file (in the same format) named 'squared-enhanced-timings'
- With gnuplot we can graph both timings:
plot [:][:] "single-timings" using 1:2 with line,
"squared-enhanced-timings" using 1:2 with line
- You can append more and more data files in this manner.




Summary

- Hopefully you can see from today that we can:
 - Analyze algorithm analytically to predict performance.
 - Profile code to find what piece of code is the bottleneck.
 - Get & plot timings to see actual performance.
- Performance Analysis could take the whole class time.
- We stick with the basics for this class.
- I do want to alert you to a couple of things that will help!

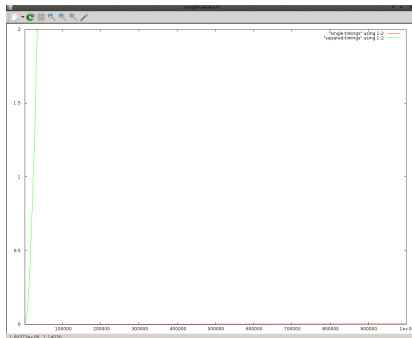
Gotcha 1

- You plot the data of two algorithms, but you can see only one!



Gotcha 1

- You plot the data of two algorithms, but you can see only one!
- Check your data and axis, usually it is because it is too small to see.




Gotcha 1

- You plot the data of two algorithms, but you can see only one!
- Check your data and axis, usually it is because it is too small to see.

- `plot [:][:] "single-timings" using 1:2 with line,`
`"squared-timings" using 1:2 with line ->`

`plot [:][:2] "single-timings" using 1:2 with line,`
`"squared-timings" using 1:2 with line`



Performance Tips

- Make sure you are working on optimizing the correct function and looking for the correct code improvements.
- From our example:

Samples: 991 of event 'cycles:ppp', Event count (approx.): 2061166086

| Overhead | Command | Shared Object | Symbol |
|----------|---------|-------------------|---------------------------|
| 95.29% | a.out | a.out | [.] sumOfOneTo |
| 4.39% | a.out | ld-2.28.so | [.] 0x0000000000000a25a |
| 0.20% | a.out | a.out | [.] sumOfOneToSquared |
| 0.08% | a.out | [kernel.kallsyms] | [k] __do_page_fault |
| 0.03% | a.out | [kernel.kallsyms] | [k] kmem_cache_alloc |
| 0.00% | perf_4. | [kernel.kallsyms] | [k] native_sched_clock |
| 0.00% | perf_4. | [kernel.kallsyms] | [k] native_apic_mem_write |
| 0.00% | perf_4. | [kernel.kallsyms] | [k] native_write_msr |

- Most of the time is spent in sumOfOneTo, so improving sumOfOneTo's performance may help.
- Improving performance of sumOfOneToSquared would not help that much.
 - Maybe we can memoize past results to use in the future – or use a better data structure.

Summary

- Again, there is a lot and we are scratching the surface.
- Important outcomes:
 - Be able to analytically deduce the performance of code.
 - Be able to profile code to find the hot spots.
 - Be able to empirically run programs to evaluate performance.
 - Understand there are anomalies that will not be addressed in this class.